

Evidence before confidence.

A practical design for turning scattered work into evidence-backed decisions, reports, applications, and public records with Kimi.

KIMI K2.5

EVIDENCE GRAPH

HUMAN RESOLUTION

CODE VALIDATION

01 / MAP THE EVIDENCE

facts / conflicts / unknowns / risks



02 / BUILD INSIDE THE BOUNDARY

cited deliverable / warnings / recipe



03 / VALIDATE WITHOUT ANOTHER MODEL

source IDs / structure / ownership / bounds

A common problem hiding in different workflows

People rarely lack work. They lack one reliable place that explains what happened, which record is final, and what can be claimed without exaggeration.

1 Student builder

Repositories, demos, certificates, event work and notes are disconnected.

Verifiable portfolio case study

2 NGO / community

Registration, check-ins, photos, surveys and reports often disagree.

Partner or donor impact report

3 Founder

Customer calls, experiments, market research and decisions lose their context.

Evidence-backed decision memo

4 Creator / researcher

Large source sets make weak claims easy to repeat and hard to trace.


Cited brief or article outline

The product promise

- Show what the supplied evidence supports.
- Expose conflicts, uncertainty and missing proof.
- Ask a human before resolving material disagreement.
- Create useful outputs with traceable source IDs.
- Validate mechanical guarantees in ordinary code.

From raw material to a defensible output

The model does two different jobs. A human resolves material conflicts. Code checks the guarantees that do not require model judgment.



What the human sees

97 appears in the organizer record. 143 appears in a draft caption. Which figure should be approved for the final report?

What code can guarantee


Every cited source ID exists. Required fields are present. Confidence is bounded. The saved result belongs to the authenticated account.

Boundary: ProofGraph does not claim to prove universal truth.

It makes the supplied evidence, conflicts, decisions and validation steps visible enough to audit.

A small architecture with explicit responsibility

Each layer owns a different kind of trust. The model is powerful, but it is not responsible for authentication, cost control or referential integrity.




Current production boundaries

Server-side provider secret | authenticated runs | account ownership | source and block caps | per-user and global quotas | same-origin writes

Start narrow, learn from corrections, then scale

The current build proves the workflow. The roadmap adds provenance, human resolution and accessible capture before expanding into organization-scale infrastructure.



Measure product quality

Conflicts surfaced / unsupported claims blocked / human corrections reused / time to approved output

Measure adoption

First evidence pack / repeat workflows / exported outputs / multilingual completion

A convincing response is an output. An inspectable process is a product.

Version zero is already running in public

ProofPack and Kimi Workflow Lab demonstrate the technical core today. The results below describe one bounded synthetic test, not a general reliability claim.

5

source blocks

2

Kimi model stages

5/5

code checks passed

27.055 s

recorded live latency

Observed in the disposable run

- Selected the organizer record of 97 over an unverified draft of 143.
- Flagged the instruction to invent 10,000 attendees and an award.
- Preserved measurable follow-up: eight repositories and three session requests.
- Used only source IDs that existed in the submitted evidence pack.

What this does not prove

The sources were synthetic. The validator checks references and structure, not external factual truth. A permission-safe evaluation set is the next requirement.

Inspect the working system

Live lab: proofpack-kimi-arun.arunchandel1780.workers.dev/lab | Source: github.com/Arun5768/proofpack-kimi